

```
In [3]: #importing needed libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import math
import seaborn as sns
```

```
In [17]: df = pd.read_csv("Boston.csv")
```

```
In [18]: df.head()
```

Out[18]:

	Unnamed: 0	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black	lstat	medv
0	1	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
1	2	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
2	3	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
3	4	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
4	5	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.2

```
In [6]: #Checking for null values
df.isnull().sum()
```

```
Out[6]: Unnamed: 0      0
      crim      0
      zn      0
      indus      0
      chas      0
      nox      0
      rm      0
      age      0
      dis      0
      rad      0
      tax      0
      ptratio      0
      black      0
      lstat      0
      medv      0
      dtype: int64
```

Q4) Use the boston.csv dataset and determine the best 5 features to predict 'MEDV'


```
In [7]: #Checking the correlation of other variables with MEDV
df.corr().tail(1)
```

```
Out[7]:
```

	Unnamed: 0	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black	lstat
medv	-0.226604	-0.388305	0.360445	-0.483725	0.17526	-0.427321	0.69536	-0.376955	0.249929	-0.381626	-0.468536	-0.507787	0.333461	-0.737

In [8]: *#Plotting correlation heatmap for a better understanding of interdependence of variables*

```
plt.figure(figsize=(16,10))
dataplot = sns.heatmap(df.corr(), cmap="YlGnBu", annot=True)
plt.show()
```



Top correlated features with MEDV are:

1. LSTAT(-0.74)
2. RM(0.7)
3. PTRATIO(-0.51)
4. INDUS(-0.48)
5. TAX(-0.47)
6. NOX(-0.43)
7. CRIM(-0.39)

Here, (INDUS-TAX), (INDUS-NOX) are highly correlated among themselves, hence both can't be selected together

So we can take up the next correlated variable in order, 'CRIM'


Q5) Using `sklearn.linear_model`, find the multiple regression model for the `boston.csv` dataset using the best 4 features. (from `sklearn.linear_model` import `LinearRegression`)

```
In [9]: # Top 4 features selected - LSTAT, RM, PTRATIO, INDUS
```

```
#Plotting these features for a visual ally
```

```
plt.figure(figsize=(20, 5))  
features = ['lstat', 'rm', 'ptratio', 'indus']
```

```
for i, col in enumerate(features):  
    plt.subplot(1, len(features) , i+1)  
    x = df[col]  
    y = df['medv']  
    plt.scatter(x, y, color = 'r')  
    plt.title(col)  
    plt.xlabel(col)  
    plt.ylabel('MEDV')
```



From the above plots, we can infer that LSTAT and RM have a few outliers but mostly follow a linear distribution, on the other hand, PTRATIO and INDUS have quite a scattered composition.

In [10]: *#Alloting values for x and y respectively*

```
X = pd.DataFrame(np.c_[df['lstat'], df['rm'], df['ptratio'],df['indus']], columns = ['lstat','rm', 'ptratio', 'rm']) #True  
Y = df['medv']
```

In [11]: `print(X.head())`
`print("\n")`
`print(y.head())`

	lstat	rm	ptratio	rm
0	4.98	6.575	15.3	2.31
1	9.14	6.421	17.8	7.07
2	4.03	7.185	17.8	7.07
3	2.94	6.998	18.7	2.18
4	5.33	7.147	18.7	2.18

0	24.0
1	21.6
2	34.7
3	33.4
4	36.2

Name: medv, dtype: float64

```
In [12]: from sklearn.model_selection import train_test_split

# Splitting the dataset into training and testing sets
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2, random_state=42)

print(X_train.shape)
print(X_test.shape)
print(Y_train.shape)
print(Y_test.shape)

(404, 4)
(102, 4)
(404,)
(102,)
```

```
In [13]: #Implementing the linear regression model for the training dataset from the sklearn library
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

lin_model = LinearRegression()
lin_model.fit(X_train, Y_train)
```

Out[13]: LinearRegression()

```
In [14]: # Assigning values to the coefficients and intercept
B0 = lin_model.intercept_
B1 = lin_model.coef_[0]
B2 = lin_model.coef_[1]
B3 = lin_model.coef_[2]
B4 = lin_model.coef_[3]
```

```
In [15]: #Printing the regression equation

print('The regression model is: y = {} + {} x1 + {} x2 + {} x3 + {} x4'
      .format(round(B0,4), round(B1, 4), round(B2, 4), round(B3, 4), round(B4, 4)))
```

The regression model is: $y = 14.663 + -0.5809 x_1 + 4.9376 x_2 + -0.8791 x_3 + 0.0262 x_4$

Q6) Find the accuracy of the model using appropriate metrics using 80, 20 split for training and test.

```
In [16]: #Importing required functions
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score

#Getting the y predicted values
y_train_predicted = lin_model.predict(X_train)
y_test_predicted = lin_model.predict(X_test)

#Defining r2 and rmse
r2 = r2_score(Y_train, y_train_predicted)
rmse = (np.sqrt(mean_squared_error(Y_test, y_test_predicted)))

#Printing the values
print('R2 score is', r2)
print('RMSE is', rmse)
```

R2 score is 0.6866881152922459

RMSE is 5.220087893560882