

Threads





Threads

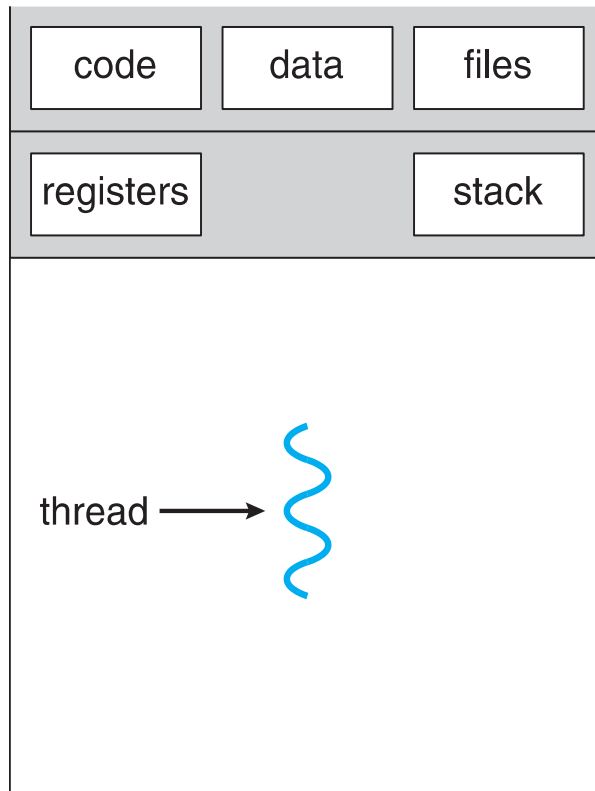
- A thread is a basic unit of CPU utilization;
 - it comprises
 - ▶ a **thread ID**,
 - ▶ a **program counter**,
 - ▶ a **register set**, and
 - ▶ a **stack**.
 - It shares with other threads belonging to the same process its **code section, data section, and other operating-system resources**, such as open files and signals.



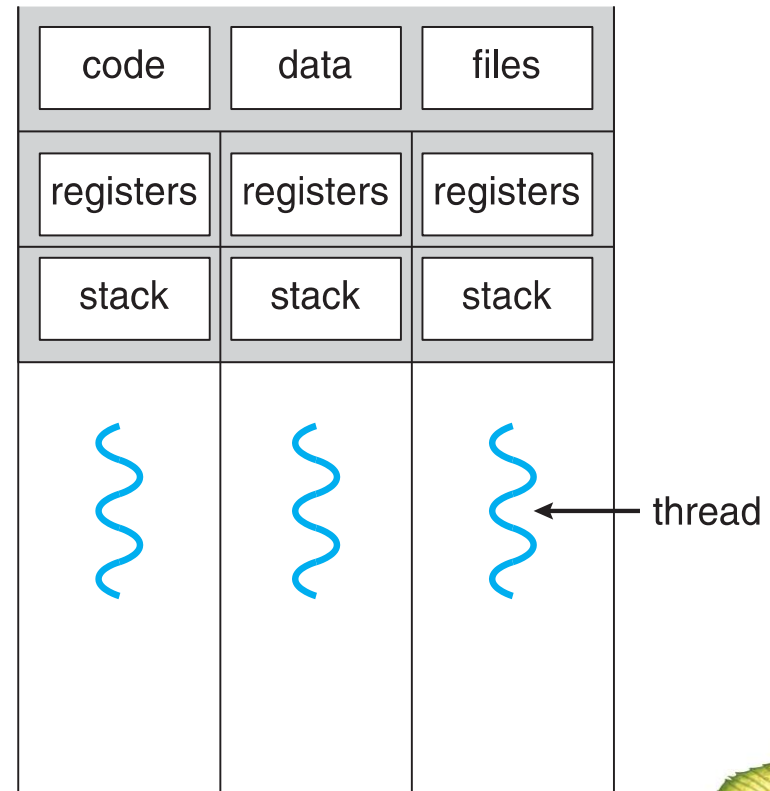


Single and Multithreaded Processes

- ✓ A traditional (or **heavy weight**) process has a single thread of control.
- ✓ If a process has multiple threads of control, it can perform more than one task at a time.



single-threaded process



multithreaded process





Motivation

- ❑ Most modern applications are multithreaded
- ❑ Threads run within application
- ❑ Multiple tasks with the application can be implemented by separate threads
 - ❑ Update display-a thread for displaying graphics
 - ❑ Fetch data-another thread for responding to keystrokes from the user
 - ❑ Spell checking-a third thread for performing spelling and grammar checking in the background
- ❑ Applications can also be designed to leverage processing capabilities on multicore systems.
- ❑ Process creation is heavy-weight while thread creation is light-weight
- ❑ Can simplify code, increase efficiency
- ❑ Kernels are generally multithreaded





Simple Thread Program

```
#include <pthread.h>
#include <stdio.h>


void* thread_code( void * param )
{
    printf( "In thread code\n" );
}

int main()
{
    pthread_t thread;
    pthread_create(&thread, 0, &thread_code, 0 );
    printf("In main thread\n" );
}
```





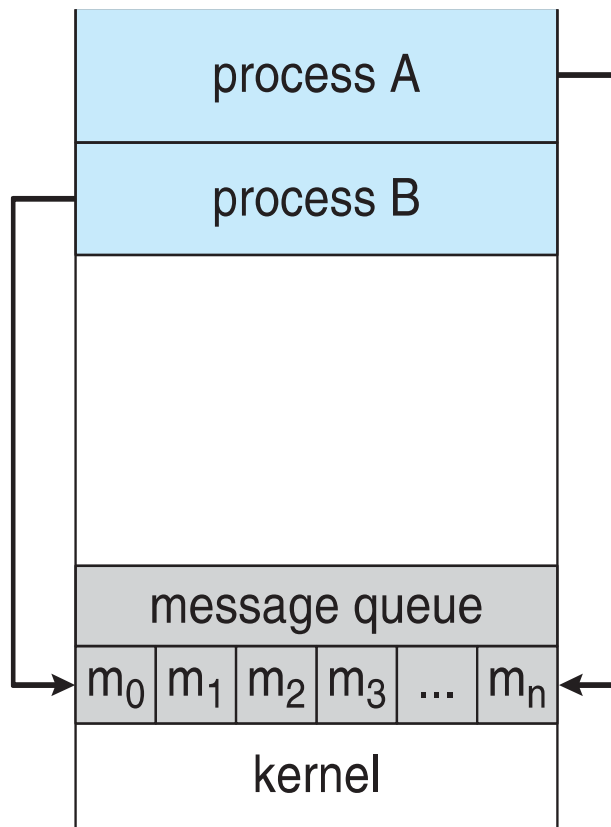
Multithreaded Server Architecture






Benefits

- **Responsiveness** – may allow continued execution if part of process is blocked, especially important for user interfaces
- **Resource Sharing** – threads share resources of process, easier than shared memory or message passing



(a)




(b)





Benefits

- ❑ **Economy** – cheaper than process creation, thread switching lower overhead than context switching
- ❑ **Scalability** – process can take advantage of multiprocessor architectures





User Threads and Kernel Threads

- ❑ **User threads** - management done by user-level threads library
 - ❑ managed without kernel support
- ❑ Three primary thread libraries:
 - ❑ POSIX **Pthreads**
 - ❑ Windows threads
 - ❑ Java threads
- ❑ **Kernel threads** - Supported by the Kernel
 - ❑ managed directly by the operating system
- ❑ Examples – virtually all general purpose operating systems, including:
 - ❑ Windows
 - ❑ Solaris
 - ❑ Linux
 - ❑ Tru64 UNIX
 - ❑ Mac OS X
- ❑ **A relationship must exist between user threads and kernel threads.**





Multithreading Models


- Many-to-One
- One-to-One
- Many-to-Many





Many-to-One


- ❑ Many user-level threads mapped to single kernel thread
- ❑ Thread management is done by the thread library in user space, **so it is efficient**
- ❑ **One thread blocking causes all to block**
- ❑ **Multiple threads may not run in parallel on multicore system because only one may be in kernel at a time**
- ❑ Few systems currently use this model
- ❑ Examples:
 - ❑ **Solaris Green Threads**
 - ❑ **GNU Portable Threads**
- ❑ very few systems continue to use the model because of its **inability to take advantage of multiple processing cores**





One-to-One


- Each user-level thread maps to a kernel thread
- Creating a user-level thread creates a kernel thread
- **More concurrency than many-to-one:** by allowing another thread to run when a thread makes a blocking system call
- It also allows multiple threads to run in parallel on multiprocessors
- Number of threads per process sometimes restricted due to overhead
- Examples
 - Windows
 - Linux
 - Solaris 9 and later





Many-to-Many Model

- ❑ Allows many user level threads to be mapped to many kernel threads (a smaller or equal number of kernel threads)
- ❑ Allows the operating system to create a sufficient number of kernel threads
- ❑ Solaris prior to version 9
- ❑ Windows with the *ThreadFiber* package






Effect of Multithreaded Models on Concurrency

- The many-to-one model allows the developer to create as many user threads as he/she wishes, it does not result in true concurrency, because **the kernel can schedule only one thread at a time.**
- The one-to-one model allows greater concurrency, but the developer has to be careful not to create too many threads within an application.
- The many-to-many model suffers from neither of these shortcomings: developers can create as many user threads as necessary, and the corresponding kernel threads can run in parallel on a multiprocessor. Also, when a thread performs a blocking system call, the kernel can schedule another thread for execution.





```
Microsoft Windows [Version 10.0.19044.1889]  
(c) Microsoft Corporation. All rights reserved.
```

```
C:\Users\MAHE>wmic
```

```
wmic:root\cli>CPU Get NumberOfCores
```

```
NumberOfCores
```

```
4
```

```
wmic:root\cli>CPU Get NumberOfCores, NumberOfLogicalProcessors
```

```
NumberOfCores    NumberOfLogicalProcessors
```

```
4
```

```
8
```

```
wmic:root\cli>
```





Thread Libraries

- **Thread library** provides programmer with API for creating and managing threads
- Two primary ways of implementing a thread libraries
 - Library entirely in user space with no kernel support
 - ▶ All code and data structures for the library exist in user space.
 - ▶ This means that invoking a function in the library results in a local function call in user space and not a system call
 - Kernel-level library supported by the OS
 - ▶ Code and data structures for the library exist in kernel space. Invoking a function in the API for the library typically results in a system call to the kernel
- Three main thread libraries are in use today: POSIX Pthreads, Windows, and Java.





Pthreads

- ❑ May be provided either as user-level or kernel-level
- ❑ A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- ❑ ***Specification***, not ***implementation***
- ❑ API specifies behavior of the thread library, implementation is up to development of the library
- ❑ Common in UNIX operating systems (Solaris, Linux, Mac OS X)





Windows, and Java Thread

- ❑ Windows thread library is a kernel-level library available on Windows systems.
- ❑ The Java thread API allows threads to be created and managed directly in Java programs.
- ❑ However, because in most instances the JVM is running on top of a host operating system, the Java thread API is generally implemented using a thread library available on the host system.
- ❑ This means that on Windows systems, Java threads are typically implemented using the Windows API;
- ❑ UNIX and Linux systems often use Pthreads.





Threading Issues

- ❑ Some of the issues to consider in designing multithreaded programs
- ❑ Semantics of **fork()** and **exec()** system calls
- ❑ Signal handling
 - ❑ Synchronous and asynchronous
- ❑ Thread cancellation of target thread
 - ❑ Asynchronous or deferred
- ❑ Thread-local storage
- ❑ Scheduler Activations





Semantics of `fork()` and `exec()`

- The semantics of the `fork()` and `exec()` system calls change in a multithreaded program
- Issue
 - If one thread in a program calls `fork()`, does the new process duplicate all threads, or is the new process single-threaded?
 - ▶ Does `fork()` duplicate only the calling thread or all threads?
- Solution
 - Some UNIX systems have chosen to have two versions of `fork()`,
 - ▶ one that duplicates all threads and
 - ▶ another that duplicates only the thread that invoked the `fork()` system call.
- But which version of `fork()` to use and when?





Semantics of `fork()` and `exec()`

- If a thread invokes the `exec()` system call, the program specified in the parameter to `exec()` will replace the entire process—including all threads.
- Depends on the application
 - If `exec()` is called immediately after forking,
 - ▶ Then duplicating all threads is unnecessary, as the program specified in the parameters to `exec()` will replace the process. In this instance, duplicating only the calling thread is appropriate.
 - If the separate process does not call `exec()` after forking,
 - ▶ Then separate process should duplicate all threads.





Signal Handling

Signals are used in UNIX systems **to notify a process that a particular event has occurred.**

- ❑ A signal may be received either **synchronously or asynchronously** depending on the **source** of and the **reason for the event** being signaled.
- ❑ All signals, whether synchronous or asynchronous, follow the same pattern:
- ❑ A **signal handler** is used to process signals
 1. Signal is generated by particular event
 2. Signal is delivered to a process
 3. Once delivered, the signal must be handled
 4. Signal is handled by one of two signal handlers:
 1. default
 2. user-defined





Signal Handling (Cont.)

Every signal has **default handler** that kernel runs when handling that signal

- ❑ **User-defined signal handler** can override default action
- ❑ Signals are handled in different ways.
 - ❑ Some signals (such as changing the size of a window) are simply **ignored**; others (such as an illegal memory access) are handled by **terminating the program**.
 - ❑ Handling signals in single-threaded programs is straightforward: signals are always delivered to a process. However, delivering signals is more complicated in multithreaded programs, where a process may have several threads. **Where, then, should a signal be delivered?**
 - ❑ Deliver the signal to the thread to which the signal applies
 - ❑ Deliver the signal to every thread in the process
 - ❑ Deliver the signal to certain threads in the process
 - ❑ Assign a specific thread to receive all signals for the process





Signal Handling (Cont.)

- The **method for delivering** a signal depends on the **type of signal generated**.
 - Synchronous signals need to be delivered to the thread causing the signal and not to other threads in the process.
 - Asynchronous signals is not as clear. Some asynchronous signals—such as a signal that terminates a process (<control><C>, for example)—should be sent to all threads.
- The standard UNIX function for delivering a signal is
 - `kill(pid_t pid, int signal)`
 - `Pthread_kill(pthread_t tid, int signal)`





Thread Cancellation

- ❑ Terminating a thread before it has finished
 - ❑ For example, if multiple threads are concurrently searching through a database and one thread returns the result, the remaining threads might be canceled.
 - ❑ Another situation might occur when a user presses a button on a web browser that stops a web page from loading any further. Often, a web page loads using several threads—each image is loaded in a separate thread. When a user presses the stop button on the browser, all threads loading the page are canceled.
- ❑ Thread to be canceled is **target thread**
- ❑ Two general approaches:
 - ❑ **Asynchronous cancellation:** One thread terminates the target thread immediately
 - ❑ **Deferred cancellation:** allows the target thread to periodically check if it should be cancelled





Thread Cancellation

- ❑ The difficulty with cancellation occurs in situations where **resources have been allocated to a canceled thread** or where a thread is canceled while **in the midst of updating data it is sharing with other threads**. This becomes especially troublesome with asynchronous cancellation. Often, the operating system will reclaim system resources from a canceled thread but will not reclaim all resources. Therefore, canceling a thread asynchronously **may not free a necessary system-wide resource**.
- ❑ With deferred cancellation, in contrast, one thread indicates that a target thread is to be canceled, but cancellation occurs only after the target thread has checked a flag to determine whether or not it should be canceled. The thread can perform this check at a point at which it can be canceled safely.
- ❑ Pthread code to create and cancel a thread:

```
pthread_t tid;

/* create the thread */
pthread_create(&tid, 0, worker, NULL);

. . .

/* cancel the thread */
pthread_cancel(tid);
```





Thread Cancellation (Cont.)

- Invoking thread cancellation requests, cancellation, but actual cancellation depends on thread state
- Pthreads supports three cancellation modes. Each mode is defined as a state and a type

Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

- If thread has cancellation disabled, cancellation remains pending until thread enables it
- Default type is deferred
 - Cancellation only occurs when thread reaches **cancellation point**
 - ▶ i.e. `pthread_testcancel()`
 - ▶ Then **cleanup handler** is invoked
- On Linux systems, thread cancellation is handled through signals





Thread-Local Storage


- ❑ **Thread-local storage (TLS)** allows each thread to have its own copy of data
- ❑ Useful when you do not have control over the thread creation process (i.e., when using a thread pool)
- ❑ Different from local variables
 - ❑ Local variables visible only during single function invocation
 - ❑ TLS visible across function invocations
- ❑ Similar to **static** data
 - ❑ TLS is unique to each thread





Scheduler Activations

- Both M:M and Two-level models require communication to maintain the appropriate number of kernel threads allocated to the application
- Typically use an intermediate data structure between user and kernel threads – **lightweight process (LWP)**
 - Appears to be a virtual processor on which process can schedule user thread to run
 - Each LWP attached to kernel thread
 - How many LWPs to create?
- Scheduler activations provide **upcalls** - a communication mechanism from the kernel to the **upcall handler** in the thread library
- This communication allows an application to maintain the correct number kernel threads





Operating System Examples

- Windows Threads
- Linux Threads





Linux Threads

- ❑ Linux refers to them as **tasks** rather than **threads**
- ❑ Thread creation is done through `clone()` system call
- ❑ `clone()` allows a child task to share the address space of the parent task (process)
 - ❑ Flags control behavior

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

- ❑ `struct task_struct` points to process data structures (shared or unique)

