

LAB-6 Programs on Strings in CUDA

In a multithreaded scenario, the issue of data inconsistency will arise, if multiple threads modify a single shared memory variable. To overcome this atomic functions need to be used. List of atomic functions, their syntax and explanation is provided below.

atomicAdd():

```
int atomicAdd (int* address, int val);  
  
unsigned int atomicAdd(unsigned int* address, unsigned int val);  
  
float atomicAdd(float* address, float val);  
  
double atomicAdd(double* address, double val);
```

Reads the 16-bit, 32-bit or 64-bit word old located at the address address in global or shared memory, computes (old + val), and stores the result back to memory at the same address. These three operations are performed in one atomic transaction. The function returns old.

Solved Example:

A CUDA program which takes a string as input and determines the number of occurrences of a character 'a' in the string. This program uses atomicAdd() function.

```
#include "cuda_runtime.h"
#include "device_launch_parameters.h"

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <conio.h>
#define N 1024

__global__ void CUDACount(char* A, unsigned int *d_count){
    int i = threadIdx.x;
    if(A[i]=='a')
        atomicAdd(d_count,1);
}

int main() {
    char A[N];
    char *d_A;
    unsigned int *count=0,*d_count,*result;
    printf("Enter a string");
    gets(A);
    cudaEvent_t start, stop;
    cudaEventCreate(&start);
    cudaEventCreate(&stop);
    cudaEventRecord(start, 0);
    cudaMalloc((void**)&d_A, strlen(A)*sizeof(char));
    cudaMalloc((void **)&d_count,sizeof(unsigned int));
    cudaMemcpy(d_A, A, strlen(A)*sizeof(char), cudaMemcpyHostToDevice);
    cudaMemcpy(d_count,count,sizeof(unsigned int),cudaMemcpyHostToDevice);
```

```

    cudaError_t error = cudaGetLastError();
    if (error != cudaSuccess) {
        printf("CUDA Error1: %s\n", cudaGetErrorString(error));
    }
    CUDACount<<<1,strlen(A)>>>(d_A,d_count);
    error = cudaGetLastError();
    if (error != cudaSuccess) {
        printf("CUDA Error2: %s\n", cudaGetErrorString(error));
    }
    cudaEventRecord(stop, 0);
    cudaEventSynchronize(stop);
    float elapsedTime;
    cudaEventElapsedTime(&elapsedTime, start, stop);
    cudaMemcpy(result, d_count, sizeof(unsigned int), cudaMemcpyDeviceToHost);
    printf("Total occurrences of a=%u",result);
    printf("Time Taken=%f",elapsedTime);
    cudaFree(d_A);
    cudaFree(d_count);
    printf("\n");
    getch();
    return 0;
}

```

Lab Exercises:

1. Write a program in CUDA to count the number of times a given word is repeated in a sentence.
(Use Atomic function)
2. Write a CUDA program that reads a string *S* and produces the string *RS* as follows:
Input string *S*: PCAP Output string *RS*: PCAPPCAPCP
Note: Each work item copies required number of characters from *S* in *RS*.

Additional Exercises:

- 1) Write a CUDA program which reads a string consisting of N words and reverse each word of it in parallel.
- 2) Write a CUDA program that takes a string *Sin* as input and one integer value N and produces an output string, *Sout*, in parallel by concatenating input string *Sin*, N times as shown below.
Input: *Sin* = "Hello" N = 3
Output: *Sout* = "HelloHelloHello"
- Note: Every thread copies the same character from the Input string *S*, N times to the required position.**
- 3) Write a CUDA program which reads a string *Sin* and produces an output string *T* as shown below.
Input: *Sin*: "Hai"
Output: *T*: "Haaiii "

Note: Every thread stores a character from input string *Sin*, required number of times into output string *T*.